

# The ECho Event Delivery System

Greg Eisenhauer

eisen@cc.gatech.edu

College of Computing  
Georgia Institute of Technology  
Atlanta, Georgia 30332

February 27, 1999 – DataExchange Version 3.1 – ECho Version 1.0

## 1 Introduction

**ECho** is an event delivery middleware system developed at Georgia Tech. Superficially, the semantics and organization of structures in ECho are similar to the Event Channels described by the CORBA Event Services specification[6]. Like most event systems, what ECho implements can be viewed as an anonymous group communication mechanism. In contrast with more familiar one-to-one send-receive communication models, data senders in anonymous group communication are unaware of the number or identity of data receivers. Instead, message sends are delivered to receivers according to the rules of the communication mechanism. In this case, event channels provide the mechanism for matching senders and receivers. Messages (or *events*) are sent via *sources* into *channels* which may have zero or more *subscribers* (or *sinks*). The locations of the sinks, which may be on the same machine or process as the sender, or anywhere else in the network, are immaterial to the sender. A program or system may create or use multiple event channels, and each subscriber receives only the messages sent to the channel to which it is subscribed. The network traffic for multiple channels is multiplexed over shared communications links, and channels themselves impose relatively low overhead. Instead of doing explicit *read()* operations, sink subscribers specify a subroutine (or handler) to be run whenever a message arrives. In this sense, event delivery is asynchronous and passive for the application.

Figure 1 depicts a set of processes communicating with event channels. The event channels are shown as existing in the space between processes, but in practice they are distributed entities, with bookkeeping data in each process where they are referenced. Channels are *created* once by some process, and *opened* anywhere else they are used. The process which creates the event channel is distinguished in that it is the contact point for other processes wishing to use the channel. The channel ID, which must be used to open the channel, contains the hostname and IP port number of the creating process (as well as information identifying the specific channel). However, event distribution is not centralized and there are no distinguished processes during event propagation. Event messages are always sent directly from an event source to all subscribers. Because of this setup, if the process which created an event channel exits the channel can no longer be opened by other processes. However, the channel will continue to function for event distribution between the process that have already opened it.

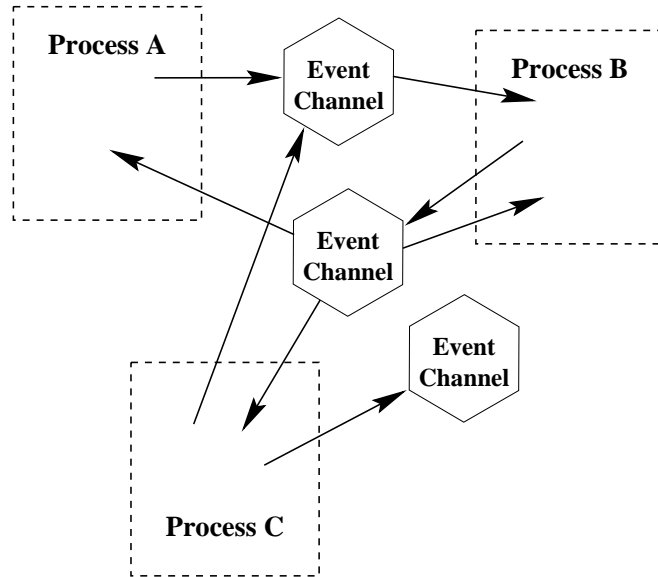


Figure 1: Processes using Event Channels for Communication.

## 2 ECho API

ECho is largely built on top of DataExchange[3] facilities. The next sections provide the basic ECho API and discuss their implementation.

### 2.1 Channel Creation and Subscription

ECho Event Channels are created with the `EChannel_create()` call given below:

```
EChannel EChannel_create(DEExchange de);
```

Channel creation essentially involves initializing a data structure that will record information about the subscribers to the event channel. The process in which the creation is performed is the permanent contact point for the event channel and participates in the subscription and unsubscription process. However, it does not necessarily participate in event propagation.

When they are created, channels are assigned a character-string global ID. This string can be used to open the channel for use in other processes. The calls involved are given below:

```
char *ECglobal_id(EChannel chan);
EChannel EChannel_open(DEExchange de, char *global_id);
```

Opening a channel that was created in another process involves obtaining a DataExchange connection to the creating process and obtaining from that process a list of other processes which have opened the channel. DataExchange connections are then obtained to each process in this list. These direct connections will be used for event delivery.

### 2.2 Event Submission

Several routines manage event submission in the ECho API. The first of these is `ECsource_subscribe()`, which derives an `ECSourceHandle` to which an event can be submitted. In the current implemen-

tation, `ECsource_subscribe` performs no significant processing, but it is a placeholder for functionality required for future features (such as the derived event channels described in Section 5). Given an `ECSourceHandle`, events can be submitted to the channel with `ECsubmit_event()` given below:

```
extern ECSourceHandle ECsource_subscribe ( EChannel chan );
extern void ECsubmit_event ( ECSourceHandle handle, void * event, int event_length );
```

Event submission is synchronous in the sense that before the submission call returns, any local handlers for the event have been called, the event has been transmitted to any remote sinks and the application is free to destroy the data that was submitted.

## 2.3 Event Delivery

Events are delivered to event sinks that have registered handlers. Handler registration is accomplished through the `EChannel_subscribe_handler()` function:

```
typedef void (*EHandlerFunction) (void *event, int length, void *client_data);
extern ECSinkHandle ECSink_subscribe( EChannel chan, EHandlerFunction func, void *client_data);
```

When an event arrives from the network or is submitted to a channel that has a local sink, the event is queued for later delivery. Event delivery consists of calling the specified handler function and passing the event as a parameter. The `client_data` value specified in the subscribe call is also passed to the handler without interpretation by the library.

Event delivery is done at two points, just before `ECsubmit_event()` returns and at the completion of a `DExchange_poll_and_handle()` operation. In each case, all pending events are dispatched. During dispatch, handler functions are called directly. Therefore, handler functions that can potentially block or compute for long periods of time should fork a thread to perform those tasks. Preservation of the event parameter data is the responsibility of the handler.

## 2.4 Simple Example Programs

Figures 2, 4 and 3 give simple independent programs for creating an event channel, submitting events to a channel and establishing a handler for events in a channel. The programs are similar in that they all involve a setup phase where they create and setup a `DataExchange` and a loop phase where they handle network traffic (via some variant of `DExchange_poll_and_handle()`). During the setup phase, each program calls `DExchange_listen()` so that it can accept connections from other programs. Generally, the use of event channels means that other programs must be allowed to initiate connections, so this listen is necessary. The channel creation program in Figure 2 is the simplest. Between the `DataExchange` setup and loop phases, all it does is create an event channel with `EChannel_create()` and print out it's ID as given by `ECglobal_id()`. This ID is used by the event source and sink programs to access the channel. This example sends a long integer which is initialized based on the time and incremented with each event submission.

```
int main()
{ /* this program creates an event channel */
  DExchange de = DExchange_create();
  EChannel chan;

  DExchange_listen(de, 0);
  chan = EChannel_create(de);
  printf("Channel ID is: %s\n", ECglobal_id(chan));
  while (1) DExchange_poll_and_handle(de, 1);
}
```

Figure 2: Simple channel creation program.

The event sending (or source) program is slightly more complex than the creation program. In addition to creating a DataExchange and doing a `DExchange_listen()`, it uses `EChannel_open()` to open a connection to the channel created by the program above. The second argument to `EChannel_open()` is the string channel ID that is printed out when the channel creation runs.<sup>1</sup> Once the channel is opened, the program identifies itself as an event source for that channel (and gets a source handle) by calling `ECsource_subscribe()`. After this, the sample program enters a loop where it submits events and uses `DExchange_poll_and_handle_timeout()` to handle the network and delay briefly between event submissions. The three arguments to `ECsubmit_event()` are the `ECsource_handle` from the `ECsource_subscribe()` call and the base address and length of the event data to send.

```
void main(argc, argv)
int argc;
char **argv;
{
    DExchange de = DExchange_create();
    EChannel chan;
    ECSourceHandle handle;
    long a = time(NULL) % 100;

    DExchange_listen(de, 0);
    chan = EChannel_open(de, argv[1]);
    handle = ECsource_subscribe(chan);
    while (1) {
        printf("I'm submitting %ld\n", ++a);
        ECsubmit_event(handle, &a, sizeof(a));
        DExchange_poll_and_handle_timeout(de, 1, 0);
    }
}
```

Figure 3: Simple event source program.

Figure 4 shows a simple program which receives events. Like the event source program in Figure 3, the sink program uses `EChannel_open()` to open the channel whose ID is passed in on the command line. After opening the channel, the program uses `ECsink_subscribe()` to register a subroutine to handle events in that channel. Finally it enters a loop that services the network using the blocking call `DExchange_poll_and_handle()`. When an event arrives, the event handler subroutine is called with the base address of the event data, the length of the data, and the `client_data` value which was specified as the third parameter to `ECsink_subscribe`. In this case, the `client_data` parameter is unused, but it could give application-specific information to the handler function.

```
void handler(event, length, client_data)
void *event;
int length;
void *client_data;
{
    printf("event data is %ld\n", *(long *) event);
}

void main(argc, argv)
int argc;
char **argv;
{
    DExchange de = DExchange_create();
    EChannel chan;
    DExchange_listen(de, 0);
    chan = EChannel_open(de, argv[1]);
    (void) ECSink_subscribe(chan, handler, NULL);
    while(1) DExchange_poll_and_handle(de, 1);
}
```

Figure 4: Simple event handling program.

The programs given in this section are very simple examples of channel creation, submission and handling programs, however they should be sufficient to demonstrate the basic principles of the channels. Note that you can run any number of event source or sink programs<sup>2</sup> and that each sink

<sup>1</sup>The ECho API doesn't provide any explicit mechanisms for communicating event channel IDs between programs. These simple example programs require the user to pass them in as command line arguments. Other programs might write them to files, send them in DataExchange messages, or exchange them in some other way.

<sup>2</sup>While the number of event sources and sinks is theoretically unlimited, there are practical limits. In particular, since event distribution is not centralized, each source process creates a DataExchange connection to each sink process.

will receive any event submitted by any source. Also notice that the channel provides no buffering. That is, events are distributed to the sinks that exist at the time the event is submitted by a source. Late joining sinks may miss early events. It's also worth noting that while these programs show event functionality divided into separate programs, that is only done for simplicity of presentation. In reality a single program can create any number of channels, and submit events to those or other channels and register handlers for any channel.

### 3 Sending Events in a Heterogeneous Environment

One glaring deficiency in the example programs in the previous section is that they do not work correctly in a heterogeneous environment where the representation of “long” may differ from machine to machine. The event channel system itself will work because it is based upon DataExchange mechanisms which handle representation differences. However, because the event channels treat event data as a simple block of bytes, binary event data (like the “long” value in the example) created on one machine may not be interpretable on another machine.

There are several possible approaches to this problem, including using `sprintf()/sscanf()` to encode/decode the data in ascii, or using `htonl/ntohl` to convert binary data to/from a standard byte order for transmission. However, this section will present examples using the relatively powerful facilities available in PBIO for encoding and decoding structured data. *Most new event channel programs should be written using **typed event channels**, described in Section 3.4. Typed event channels provide most of the functionality described in these sections, but without requiring the programmer to call PBIO directly.*

#### 3.1 Sending simple structured data with PBIO

The PBIO IOContext facilities are more fully explained in [2], but the example programs given in Figures 6 and 7 recast an example used in that paper and demonstrate how to use the facilities with event channels. Figure 5 at the right gives defines the data structure we'll send as event data in this example and gives the PBIO field list which describes the structure. These same definitions are used in both the source and sink programs, but differences in the source and sink machine architectures may cause the structure to be laid out quite differently. For example, on a 64-bit SGI the long field and the two pointer fields will each be 8 bytes. On a 32-bit x86 running Linux, those

```
typedef struct _dyn_rec {
    char      *string;
    long      icount;
    double    *double_array;
} dyn_rec, *dyn_rec_ptr;

IOField dyn_field_list[] = {
    {"string field", "string", sizeof(char *),
     IOOffset(dyn_rec_ptr, string)},
    {"icount", "integer", sizeof(long),
     IOOffset(dyn_rec_ptr, icount)},
    {"double_array", "float[icount]", sizeof(double),
     IOOffset(dyn_rec_ptr, double_array)},
    { NULL, NULL, 0, 0}
};
```

Figure 5: Type definition and PBIO field list for sending structured data.

fields will all be 4 bytes and the bytes are arranged in the opposite order. Given the information in the IOField lists, PBIO can pack the structure, the string and the variable-sized array of doubles for transmission. Upon receipt, the encoded event is unpacked and the data converted to the record

---

Most operating systems establish limits as to the number of simultaneous sockets or file descriptors that may be in use at one time. Programs using event channels may run into this limit when a large number of sinks are present.

```

void main(argc, argv)
int argc;
char **argv;
{
    DExchange de = DExchange_create();
    EChannel chan;
    ECSourceHandle handle;
    IOContext src_context = create_IOcontext();
    IOFormat rec_ioformat;
    dyn_rec rec;
    long a = time(NULL) % 100;

    DExchange_listen(de, 0);
    chan = EChannel_open(de, argv[1]);
    handle = ECsource_subscribe(chan);

    rec_ioformat = register_IOcontext_format("dynamic format", dyn_field_list, src_context);
    rec.string = "Hi Mom!";
    rec.icount = time(NULL) % 10 + 1;
    rec.double_array = (double*) malloc(sizeof(double) * rec.icount);

    while (1) {
        int i, buf_size;
        char *encoded_buffer;
        for (i=0; i<rec.icount; i++) rec.double_array[i] = (a + i) * 2.717;
        encoded_buffer = encode_IOcontext_buffer(src_context, rec_ioformat, &rec, &buf_size);
        printf("I'm submitting\n");
        ECsubmit_event(handle, encoded_buffer, buf_size);
        DExchange_poll_and_handle_timeout(de, 1, 0);
        a++;
    }
}

```

Figure 6: Event source program with PBIO encoding.

format in use on the receiving machine.

Figure 6 shows the modified event source program. Aside from the code which initializes the event data structure, there are three lines which relate to the PBIO encoding. The first is the declaration and initialization of a PBIO `IOContext` variable to be used in later calls. The second important change is the addition of a call to `register_IOcontext_format()`. This registers the record format with the `IOContext` and returns an `IOFormat` value that can be used to encode the record. Finally, the call to `encode_IOcontext_buffer()` does the actual encoding of the base structure, string and variable-sized double array. It returns a pointer to the encoded buffer and sets the `buf_size` to the length of that buffer. (The actual buffer used occupies temporary memory associated with the `IOContext`. The buffer will be valid until the next call to the `IOContext`.) The encoded buffer and its length then become the parameters to the `ECsubmit_event()` call.

The corresponding sink program is shown in Figure 7. This program also creates an `IOContext` variable. Because this only needs to be done once, the variable is declared `static` and initialized only the first time the handler is called. The next line of the handler uses `get_format_IOcontext()` to get the `IOFormat` handle associated with the encoded event data. Given this handle, we check to see if there is a conversion associated with that `IOFormat` and if not, we set one with

```

void handler(event, length, client_data)
void *event;
int length;
void *client_data;
{
    static IOContext dest_context = NULL;
    IOFormat dyn_rec_ioformat;
    dyn_rec rec;
    int i;

    if (dest_context == NULL) dest_context = create_IOcontext();
    dyn_rec_ioformat = get_format_IOcontext(dest_context, event);
    if (!has_conversion_IOformat(dyn_rec_ioformat)) {
        set_conversion_IOcontext(dest_context, dyn_rec_ioformat, dyn_field_list, sizeof(dyn_rec));
    }
    decode_IOcontext(dest_context, event, &rec);
    printf("event data is -> string = \"%s\"\n", rec.string);
    printf("            icount = %d\n", rec.icount);
    for (i=0; i< rec.icount; i++) {
        printf("            double_array[%d] = %g\n", i, rec.double_array[i]);
    }
}

void main(argc, argv)
int argc;
char **argv;
{
    DExchange de = DExchange_create();
    EChannel chan;

    DExchange_listen(de, 0);
    chan = EChannel_open(de, argv[1]);
    (void) ECsink_subscribe(chan, handler, NULL);
    while(1) DExchange_poll_and_handle(de, 1);
}

```

Figure 7: Event sink program with PBIO encoding.

`set_conversion_IOcontext()`. Setting a conversion for a particular record format is similar to registering a format with `register_IOcontext_format()` except that instead of telling PBIO in what record layout we'll be giving it data, it specifies the record layout we want after the record is decoded. Finally, `decode_IOcontext()` is used to convert the encoded buffer into the receiver's native record structure.

Some of the finer points of PBIO IOContexts are worth discussing at this point. One point in particular affects the number of times that `set_conversion_IOcontext()` is called. The format that is returned by `get_format_IOcontext()` represents the record format on the sending machine. In a homogeneous system, all the senders have the same record format, `get_format_IOcontext()` always returns the same value and the conversion is only set once. However, in a heterogeneous system, `get_format_IOcontext()` will return a different format for each different architecture that differs in byte order, field size or record layout. In this case, the conversion must be set for each different incoming format. This simple example uses the same native field list for each conversion, resulting in all the incoming records being converted into the same native record format. However,

a more complex program might make other choices. For example, a receiver might make different choices based on the fields that are actually present in the incoming record format, compensating for missing fields by initializing local values before doing the decode call. A program that makes intelligent decisions based on the fields actually transmitted can more easily tolerate changes in other clients. For example, newer clients may have an additional field in their events that is not present in events sent by older clients. By virtue of PBIO, the older clients will automatically ignore the new field, but the newer clients must be careful in setting their conversions. Use the new field list to set conversion on records from older clients will fail because the new field is not present. A working approach is to differentially set conversions for those formats using the old field list and handle those records specially.

### 3.2 Using nested structures

The examples in the previous section work well for relatively simple structures, but they do not properly handle nested structures. Changing the event sender to handle nested structures is easy. All that is necessary is to register the substructure formats before the structures they are used in. Consider another example borrowed from [2], the nested structures shown in Figure 8 at the right, and the field lists below:

```
static IOField R3field_list[] = {
    {"x", "float", sizeof(double), IOOffset(R3vector*, x)},
    {"y", "float", sizeof(double), IOOffset(R3vector*, y)},
    {"z", "float", sizeof(double), IOOffset(R3vector*, z)},
    {NULL, NULL, 0, 0},
};

static IOField particle_field_list[] = {
    {"loc", "R3vector", sizeof(R3vector), IOOffset(particle*, loc)},
    {"deriv1", "R3vector", sizeof(R3vector), IOOffset(particle*, deriv1)},
    {"deriv2", "R3vector", sizeof(R3vector), IOOffset(particle*, deriv2)},
    {NULL, NULL, 0, 0},
};
```

```
typedef struct R3vector_struct {
    double x, y, z;
} R3vector;

typedef struct particle_struct {
    R3vector loc;
    R3vector deriv1;
    R3vector deriv2;
} particle;
```

Figure 8: A nested record format

To send this record in the example given in Figure 6 the only PBIO-related change is to register both record formats where previously only one was registered. That is, the line containing `register_IOcontext_format()` expands to the lines:

```
(void)register_IOcontext_format("R3vector", R3field_list, src_context);
rec_ioformat = register_IOcontext_format("particle", particle_field_list, src_context);
```

Note that the IOFormat return value of the first registration is discarded. A robust program would check it for NULL to guard against failure, but otherwise it is unused.

Changes in the sink program from Figure 7 are a little more complex. The check to see if the record already has a conversion is still appropriate and will serve the same function. However, before we can do a `set_conversion_IOcontext` on the record format, we must set conversions for any subformat it may contain. For this we use the PBIO routine `get_subformats_IOcontext()` in a manner illustrated below (this code would replace the single `set_conversion_IOContext()` call inside the if statement in the handler):



```

int i = 0;
IOFormat *subformats;
subformats = get_subformats_IOcontext(dest_context, event);
while (subformats[i] != NULL) {
    char *subformat_name = name_of_IOformat(subformats[i]);
    if (strcmp(subformat_name, "R3vector") == 0) {
        set_conversion_IOcontext(dest_context, subformats[i], R3field_list, sizeof(R3vector));
    } else if (strcmp(subformat_name, "particle") == 0) {
        set_conversion_IOcontext(dest_context, subformats[i], particle_field_list,
                                sizeof(particle));
    } else {
        printf("Unexpected format in event record, \"%s\".\n", subformat_name);
    }
    i++;
}
free(subformats);

```

The code above is relatively straightforward. For each format that is used in the record, it sets the appropriate conversion based on the name associated with the format. As described in the previous section, this code could make use of more information than just the format name in determining the fields that it wanted to use, but the name is sufficient for this example.

### 3.3 More efficient encodes

One downside of the functionality of the `encode_IOcontext_buffer()` call is that it must copy all of the event data in order to create a contiguous buffer of encoded data. For small events, this isn't much of an issue, but for large events copy operations may amount to a significant fraction of network transmission overhead. To support efficient encoding of large amounts of data, PBIO has another style of data encoding whose output is not a contiguous block of encoded data. Instead its output is a list of sections of memory which together constitute the encoded data. Some of those memory sections will be in temporary memory because makes copies of structures containing pointers (PBIO must change pointers into buffer offsets before sending those structures), but elements of the event data which don't contain pointers will appear in the list directly without copying. Encoding large blocks of data with this new mechanism, `encode_IOcontext_to_vector()`, has computational costs more closely tied to the number of data blocks rather than their aggregate size. This can be a huge win, with large dynamic arrays encoded in constant time regardless of their actual size.

In order to support the use of `encode_IOcontext_to_vector()` with events there is another form of `ECsubmit_event()` that allows the event data to be passed as a list (or vector) of buffers instead of a contiguous block. With the exception of this difference, `ECsubmit_eventV()` behaves identically to `ECsubmit_event()`. Changes to the event source program given in Figure 6 to use these routines are given below:

```

IOEncodeVector encoded_buffer;
encoded_buffer = encode_IOcontext_to_vector(src_context, rec_ioformat, &rec);
ECsubmit_eventV(handle, encoded_data);

```

The type of `encoded_buffer` changes from `char *` to `IOEncodeVector`, which is really just a list of `<address, length>` pairs (NULL terminated). The `buf_size` local variable which held the length of the encoded buffer is no longer needed since the total buffer length is available from

the `IOEncodeVector`. `encode_IOcontext_to_vector()` replaces `encode_IOcontext_buffer()` and `ECsubmit_eventV()` replaces `ECsubmit_event`. The `IOEncodeVector` and some portions of the encode data are held in temporary memory associated with the `IOcontext`. They must not be freed and they will remain valid until the next `IOcontext` operation.

### 3.4 Typed Event Channels

Echo also offers typed event channels which internally provide much of the functionality discussed above. However, it also imposes some restrictions. For example, if `PBIO` is used directly, programmers have more flexibility in terms of submitting records of multiple types in the same event channel and using customized decoding and handling for each type. The typed event channel interface restricts event channels to carrying a single event type, but for that event type it handles all the conversions necessary to exchange binary data across heterogeneous machines. The remainder of this section details their use and interface.

Unlike untyped event channels, typed channels require type information to be specified at the time of channel creation. `PBIO` field lists are used to specify this information and the list of field names and types become the “type” of the event channel. (Field size and offset can also be specified in the `IOFieldList`, but those values do not become part of the channel type specification.) In order to allow specifications of the field lists associated with structured subfields in the type, a null-terminated list of subfield type names and field lists can also be specified in the `format_list` parameter.

```
typedef struct _DEformat_list {
    char *format_name;
    IOFieldList field_list;
} *DEFormat, DEFormatList;

extern EChannel
EChannel_typed_create(DEExchange de, IOFieldList field_list, DEFormatList format_list);
```

If there are no structured subfields, the `format_list` parameter should be `NULL`. Thus the code to create a typed event channel which will carry events of the form described in Figure 8 looks like this:

```
static DEFormatList particle_format_list[] = {
    {"R3vector", R3field_list},
    {NULL, NULL},
};

chan = EChannel_typed_create(de, particle_field_list, particle_format_list);
```

where `R3field_list` and `particle_field_list` are as given in Section 3.2. Please note that both the field list and the format list are declared globally. These lists must remain valid for the duration of the channel’s existence. Most commonly they will be known statically at compile time and can be declared as in this example. They should *not* be stack-allocated variables. They can be dynamically allocated, but in that case the programmer is responsible for freeing those structures after the channel is destroyed.

Once a typed channel is created remote access can be accomplished through the standard

`ECglobal_id()` and `EChannel_open()` calls. However, there are special event submission and sink and source subscribe calls as given below:

```
extern ECSourceHandle
ECsource_typed_subscribe (EChannel chan, IOFieldList field_list, DEFormatList format_list);
extern void
ECsubmit_typed_event (ECSourceHandle handle, void *event);

typedef void (*ECTypedHandlerFunction) (void *event, void *client_data);
extern ECSinkHandle
ECsink_typed_subscribe (EChannel chan, IOFieldList field_list, DEFormatList format_list,
                       ECTypedHandlerFunction func, void *client_data);
```

Essentially, these calls are the same as the corresponding untyped calls except for the addition of the `field_list` and `format_list` parameters to the subscribe functions and the elimination of the event length parameter in the submit and handler functions. The `field_list` and `format_list` specified in the subscribe calls *must* be appropriate for the type of channel they are applied to. Generally that means that the field names and types specified in the subscribe `field_list` and `format_list` must exactly match those in the `ECchannel_typed_create()` call.<sup>3</sup> Note that while the field sizes and offsets in each of these calls must exactly specify the formats of those records on the machine executing the call, there is no requirement that these match those specified for the channel or any other sink. Any necessary transformations are provided by the channel. Additionally, if a channel is created with `EChannel_typed_create()`, all source and sink subscriptions must be typed.

## 4 Events and Threads

The previous examples in this document have all assumed non-threaded execution environment. In moving to threaded environments there are a number of issues to clarify, including questions of simultaneous use of event channels, determining which thread actually executes the event handlers and thread safe use of PBIO. This section will attempt to address these issues.

### 4.1 Thread basics

All DataExchange routines, including the ECho routines, are capable of performing locking to protect their own data structures from simultaneous access. To avoid having different versions of DataExchange for each thread package which a program might use, DataExchange uses a package called `gen_threads`. `Gen_threads` is a generic wrapper for threads packages which provides an abstract thread interface implemented with function pointers. `Gen_thread` initialization calls fill in function pointers appropriate to specific thread packages so that DataExchange routines can use them. If you use DataExchange in a threaded environment, you must initialize `gen_threads` before calling `DExchange_create()` so that the appropriate locks are created. If you initialize `gen_threads` later something bad is likely to happen.

Currently there are three different sets of wrappers with `gen_threads`. `gen_pthread_init()` will initialize `gen_threads` to use the Pthreads package. Calling `gen_cthread_init()` sets up the

---

<sup>3</sup>There is some room for variance here. In particular, it is ok to provide *more* information than necessary and to ask for *less* information than is available. So the fields and formats provided for the source subscribe can be a superset of those associated with the channel. Conversely the fields and formats provided for the sink subscribe can be a subset of those of the channel. Excess information is discarded at the sink end, so it does travel over the network.

```

int    gen_thr_initialized();
thr_thread_t thr_fork(void_arg_func func, void *arg);
void   thr_thread_detach(thr_thread_t thread);
void   thr_thread_yield();
thr_mutex_t thr_mutex_alloc();
void   thr_mutex_free(thr_mutex_t m);
void   thr_mutex_lock(thr_mutex_t m);
void   thr_mutex_unlock(thr_mutex_t m);
thr_condition_t thr_condition_alloc();
void   thr_condition_free(thr_condition_t m);
void   thr_condition_wait(thr_condition_t c, thr_mutex_t m);
void   thr_condition_signal(thr_condition_t c);
void   thr_condition_broadcast(thr_condition_t c);
void   thr_thread_exit(void *status);
int    thr_thread_join(thr_thread_t t, void **status_p);
thr_thread_t thr_thread_self();

```

Figure 9: Gen\_threads API

Georgia Tech Cthreads package instead. In addition to linking with the gthreads library these routines also require you to link with the appropriate external libraries to support the threads package. Initializing gen\_threads appropriately is necessary to protect DataExchange, but you also have the option of using the gen\_threads interface in your application. Doing this instead of directly calling the threads package of your choice may make your application more portable in the sense that you may not have to rewrite your code to change threads packages. However, be warned that gen\_threads does not make all threads packages behave the same. For example, user-level and kernel-level threads packages behave differently when a thread does a blocking call. In the former case, all the threads in the process are blocked and in the latter only the calling thread is blocked. Using gen\_threads does not hide this or other important semantic differences that may exist between thread libraries. It only provides a generic API to common thread functionality. The header file in Figure 9 summarizes the interface offered by gen\_threads. The names of these calls link them with their obvious counterparts in common thread packages and should be familiar to most thread programmers. The only one which might be non-obvious is `gen_thr_initialized()` which simply returns true if gen\_threads has been initialized.

## 4.2 Threads and Event Handler Routines

Multithreaded programs which use DataExchange or ECho almost always devote a single thread to handling the network (through `DExchange_poll_and_handle()` with the `block` parameter set to `TRUE`). A less common alternative which works out essentially the same is to have a single thread poll the network at intervals (`block` parameter set to `FALSE` so `DExchange_poll_and_handle()` returns without waiting).<sup>4</sup>

Given that there may also be other application threads which may submit events, create channels and/or register event handlers, which of these threads actually runs the event handlers? The answer depends upon where the event is coming from. If the event was submitted in another process and therefore arrives over the network, then the thread which calls `DExchange_poll_and_handle()` performs the call to the handler. One side effect of this is that the network will not be serviced during the time that an event handler is running. Because of this, long running handlers may want

---

<sup>4</sup>Multiple threads calling `DExchange_poll_and_handle()` on the same `DExchange` value is likely to have unpredictable results. Don't try this at home.

to fork a thread to complete their work. In the case of a handler forking a thread, be aware that the event data is only valid until the handler returns. If the forked thread needs access to the event data it must preserve it somehow.

If a channel has event sources and sinks in the same process, then a different situation results and it is the thread which calls `ECsubmit_event()` which is borrowed to run the event handler. This situation has two side effects. First, as described in the paragraph above, long running handlers may want to fork a thread instead of delaying the submitting thread. Second, it is possible for handlers which themselves submit events to create an “event avalanche.” That is, the first call to `ECsubmit_event()` results in a call to a local handler which itself calls `ECsubmit_event()` which calls a local handler which calls `ECsubmit_event()` which calls a local handler which calls `ECsubmit_event()`... This is not *per se* a problem for the event channel system as long as the chain ends eventually (otherwise it’s an infinite recursion), but there may be other consequences with respect to stack usage and delay of the original submitting thread that may be of concern.

### 4.3 Assigning Handlers to a Specific Thread

ECho does offer a mechanism for overriding the normal way that event handlers are executed and instead assign them to a specific thread. More specifically, there is a channel subscribe function, `EChannel_subscribe_context()`, that allows event handlers to be assigned to a specific `DExchange` value. The thread performing `DExchange_poll_and_handle()` on that `DExchange` will then execute the handler. Note that there need not be any network connections associated with a `DataExchange`. In this case, the `DataExchange` can be used solely for event handler execution.

The existence of this functionality complicates event submission somewhat. Previously, the callers of `ECsubmit_event` could assume that they were free to destroy event data as soon as that routine returned. This was a side effect of local events handlers being executed by the thread calling `ECsubmit_event`. However, if event handlers can be assigned to other threads, that assumption is no longer valid and a mechanism must be established for preserving event data until the handler has an opportunity to run. That mechanism is a new event submission routine, `ECsubmit_general_event()`. This routine has an additional parameter to allow the application to specify an appropriate deallocation routine for the event data. This allows the event system to keep the event data until all local handlers have run and then deallocate it. The application should not reference the event data after it has been passed to `ECsubmit_general_event()`. *If `EChannel_subscribe_context()` is used locally with a channel, then all event submissions to that channel must be done using `ECsubmit_general_event()` instead of `ECsubmit_event()`.* Since this applies only to local handlers, there is no global requirement to know if other processes use `EChannel_subscribe_context()`, you only have to know if your program uses it.

To complement the `EChannel_subscribe_context()` and `ECsubmit_event()` calls described above, there are similar routines for typed event channels. The API for all these calls follows:

```

extern ECSinkHandle
ECsink_subscribe_context (EChannel chan, EHandlerFunction func, void *client_data,
                        DExchange de);

extern ECSinkHandle
ECsink_typed_subscribe_context (EChannel chan, IOFieldList field_list, DEFormatList format_list,
                               ECTypedHandlerFunction func, void *client_data, DExchange de);

extern void
ECsubmit_general_event (ECHandle handle, void *event, int event_len,
                       EventFreeFunction free_func);

extern void
ECsubmit_general_typed_event (ECHandle handle, void *event,
                              EventFreeFunction free_func);

extern void
ECsubmit_general_eventV (ECHandle handle, IOEncodeVector eventV,
                        EventFreeFunction free_func);

```

Figure 10 shows a program which handles events in multiple threads. This example is written using the `gen_threads` package to access the thread library and calls `gen_pthread_init()` to bind those calls to the corresponding Pthread routines. The main program creates an event channel and forks two threads, each of which subscribe to the channel and events on private `DEExchange` values. `thread1_func` blocks waiting for events, while `thread2_func` polls every 10 seconds. The main thread submits events to the channel at half second intervals. When this program is run it will print out “I’m submitting xxx” every half second. Immediately after event submission, the first thread is woken to handle the event and its handler is called. Every ten seconds, the second thread will poll for events and its handler will be called for all of its events (typically 20 events). Note that in this example, the call to `ECsink_subscribe_context()` is made by the thread which will handle that `DEExchange`, but that is not strictly necessary. The subscribe call could just as easily have been done in the main subroutine. Also note that the same event handler subroutine is used by each thread, but its output is customized using the `client_data` parameter in the subscribe call.

## 4.4 Threads and PBIO Encoding

One final topic under the subject of threads in the event system relates to the use of PBIO for event data encoding. While `DataExchange` protects itself from the hazards of concurrent use by threads, PBIO does not. Therefore, the examples of direct use of PBIO given in Sections 3.1 through 3.3 are unsafe in a threaded environment.<sup>5</sup> Besides the possibility of simultaneous format registrations threatening the consistency of the `IOcontext`, the calls `encode_IOcontext_buffer()`, `encode_IOcontext_to_vector()` and `decode_IOcontext()` leave at least portions of their results in temporary memory associated with the `IOcontext`. This memory is valid until the next `IOcontext` operation, but if multiple threads share an `IOcontext` then it may change unpredictably.

One obvious solution is to devote an `IOcontext` per thread. This is acceptable, particularly when there are few long-lived threads, but it may have excessive overhead in other circumstances. Separate `IOcontexts` do not share information, so they will each end up downloading and replicating information from the format server. To avoid this overhead, we can use a single `IOcontext`, some simple locking procedures and different PBIO calls. For example, to modify the event submission

---

<sup>5</sup>While the examples in Sections 3.1 through 3.3 are not thread-safe, this is not true for the typed event channel example in Section 3.4. `ECho` provides all necessary locking to ensure that the encoding and decoding associated with typed event channels is thread-safe.

```

static EChannel chan;

static void handler(event, length, client_data)
void *event;
int length;
void *client_data;
{
    printf("In the %s handler, thread (%lx), event is %ld\n", client_data,
           (long)thr_thread_self(), *(long *) event);
}

static void thread1_func(de)
DExchange de;
{
    ECsink_subscribe_context(chan, handler, "first", de);
    while(1) { /* block waiting for events */
        DExchange_poll_and_handle(de, 1);
    }
}

static void thread2_func(de)
DExchange de;
{
    ECsink_subscribe_context(chan, handler, "second", de);
    while(1) { /* handle events every 10 sec */
        sleep(10);
        DExchange_poll_and_handle(de, 0);
    }
}

void main()
{
    DExchange de, de1, de2;
    ECSourceHandle handle;
    long a = time(NULL) % 100;

    gen_pthread_init();
    de = DExchange_create();
    DExchange_listen(de, 0);
    chan = EChannel_create(de);
    handle = ECsource_subscribe(chan);

    de1 = DExchange_create();
    thr_fork((void_arg_func)thread1_func, (void*)de1);
    de2 = DExchange_create();
    thr_fork((void_arg_func)thread2_func, (void*)de2);

    while (1) {
        long *event = malloc(sizeof(long));
        *event = a++;
        printf("I'm submitting %ld\n", *event);
        ECsubmit_general_event(handle, event, sizeof(*event), free);
        DExchange_poll_and_handle_timeout(de, 0, 500000);
    }
}

```

Figure 10: Event handling in specific threads.

code in Figure 6 or Section 3.2 for threads, a mutex lock should be created to guard access to the `IOcontext` and lock/unlock pairs placed around the `register_IOcontext_format()` calls.

However, the situation with the temporary memory used by the encode and decode calls is more complex. We could simply acquire and hold a lock on the `IOcontext` for the entire time we need the temporary memory to remain valid. However, consider the situation of handler that is in the same process as (and sharing the `IOContext` with) the event submitter. The submitter acquires a lock on the `IOcontext` for encoding the buffer and plans to release the lock after event submission is completed (when it is done with the temporary memory). However, according to the semantics of `ECsubmit_event`, the local handler runs before `ECsubmit_event` returns. The handler will also need to acquire the lock on the `IOcontext` to protect its use of temporary memory during the decode, but the lock is already held by the submitter. This results in a deadlock and is an unworkable situation. To resolve the conflict, we have to find a way for the submitter to release the `IOContext` lock before it calls `ECsubmit_event()`.

The solution is yet another PBIO encoding routine. `encode_IOcontext_release_vector()` differs from `encode_IOcontext_to_vector()` in that it does not leave any temporary memory associated with the `IOcontext`. Instead, the portions of the encoded buffer that would have been in temporary memory are placed in a block of malloc'd memory that is thereafter "owned" by the caller. Since nothing related to the encoding is left with the `IOcontext`, protective locks can be dropped immediately and do not have to be held for the duration of the `ECsubmit_event()`. The `IOEncodeVector` return value of `encode_IOcontext_release_vector()` is at the head of the malloc'd block, followed by any other non-leaf data segments that would have been left in temporary memory. Therefore, the entire block of encoded data can be freed just by passing the `IOEncodeVector` to `free()`. The following code section shows how event submission can be accomplished in a thread-safe way with this routine:

```
IOEncodeVector encoded_data;
thr_mutex_lock(context_lock);
encoded_data = encode_IOcontext_release_vector(src_context, rec_ioformat, &rec);
thr_mutex_unlock(context_lock);
ECsubmit_eventV(handle, encoded_data);
free(encoded_data);
```

Alternatively, if `EChannel_subscribe_context()` is used locally then `ECsubmit_eventV()` is not safe and `ECsubmit_general_eventV()` must be substituted.

```
IOEncodeVector encoded_data;
thr_mutex_lock(context_lock);
encoded_data = encode_IOcontext_release_vector(src_context, rec_ioformat, &rec);
thr_mutex_unlock(context_lock);
ECsubmit_general_eventV(handle, encoded_data, free);
```

On the decoding side, a similar change is possible. The call `decode_IOcontext()` can potentially leave some decoded event data in temporary memory with the `IOContext`. In particular, the base record is placed into memory provided to `decode_IOcontext()`, but any strings or dynamic elements are left in temporary memory. If that memory was to be preserved, a lock would have to be held throughout the duration of the event handler. However, the call `decode_to_buffer_IOcontext()` decodes *all* of the event data into memory that you provide. To ensure that there is enough room to hold the entire event, the routine `this_IOrecord_length()` can be used to find the total decoded



length of the event. The body of the handler subroutine in Figure 7 could be modified to be thread-safe as follows:

```

IOFormat dyn_rec_ioformat;
dyn_rec *rec;
int i;

thr_mutex_lock(context_lock);
dyn_rec_ioformat = get_format_IOcontext(dest_context, event);
if (!has_conversion_IOformat(dyn_rec_ioformat)) {
    set_conversion_IOcontext(dest_context, dyn_rec_ioformat, dyn_field_list, sizeof(dyn_rec));
}
rec = (dyn_rec*) malloc(this_IOrecord_length(src_context, event, event_len));
decode_to_buffer_IOcontext(dest_context, event, rec);
thr_mutex_unlock(context_lock);
printf("event data is -> string = \"%s\"\n", rec.string);
printf("                    icount = %d\n", rec.icount);
for (i=0; i< rec.icount; i++) {
    printf("                    double_array[%d] = %g\n", i, rec.double_array[i]);
}
free(rec);

```

To avoid race conditions, we also move the `create_IOcontext()` to the main program and add a call there to `thr_mutex_alloc()` to allocate `context_lock`, the mutex lock used to protect the `IOcontext`.

## 4.5 ECho Summary

ECho was initially developed as a data transport mechanism to support work in application-level program monitoring and steering of high-performance parallel and distributed applications[4]. In this environment, efficiency in transporting large amounts of data is of critical concern to avoid overly perturbing application execution. Because of this, ECho was designed to take careful advantage of DataExchange and PBIO features so that data copying is minimized. Even typed event transmissions require the creation of only small amounts of header data and require no copying of application data. Also, because ECho transmits events directly to sinks, it naturally suppresses event traffic when there are no listeners.

However, as in many other situations using event-based communication, program monitoring can produce large numbers of events which may overwhelm both the listeners and the intervening networks. This can be particularly frustrating if the listener is not interested in every byte of data that it receives. Unwanted events waste network resources in carrying the data, cause unnecessary perturbation to the application sending the data, and waste compute time for the listener who has to be interrupted, read, unpack and discard events he would rather not be bothered with. Using many event channels to subdivide dataflows is an effective and low-overhead of reducing unwanted traffic because listeners can limit their sink subscriptions to event channels carrying data that they want to receive. However, effective use of this technique requires the event sender have *a priori* knowledge of the appropriate subdivisions. The technique is also much more difficult to apply when a listener's definition of "unwanted event" depends upon the event content.

ECho's *Derived Event Channels* allow sink-specified event filtering, and even event data reduction, to be applied on the source end of event transmission. Performing these calculations at the source can be a win-win situation, reducing costs for both the sender and receiver and reduc-

ing bandwidth requirements on the intervening networks. The next section describes the Derived Event Channel abstraction and the critical role of dynamic code generation in performing these calculations in an efficient way in a heterogeneous environment.

## 5 Derived Event Channels

### 5.1 General Model

Consider the situation where an event channel sink is not really interested in *every* event submitted, but only wants every  $N$ th event, or every event where a particular value in the data exceeds some threshold. Considerable network traffic could be avoided if we could somehow manage to transmit only the events of interest. One way to approach the problem is to create a new event channel and interpose an event filter as shown in Figure 11.

The event filter can be located on the same node as the event source and is a normal event sink to the original event channel and a normal source to the new, or filtered, event channel. This is a nice solution in that it does not disturb the normal function of the original event channel. However, it fails if there is more than one event source associated with the original event channel. The difficulty is that, as a normal sink, the event filter must live in some specific process. If there is more than one source subscribed to the original event channel and those sources are not co-located, as shown in Figure 12, then we still have raw events traveling over the network from Process A to Process B to be filtered.

The normal semantics of event delivery schemes do not offer an appropriate solution to the event filtering problem. Yet it is important problem to solve because of the great potential for reducing resource requirements if unwanted events can be suppressed. Our approach involves extending event channels with the concept of a *derived* event channel. Rather than explicitly creating new event channels with intervening filter objects, applications that wish to receive filtered event data create a new channel whose contents are derived from the contents of an existing channel through an application supplied derivation function,  $F$ . The event channel implementation will move the derivation function  $F$  to all event sources in the original channel, execute it locally whenever events

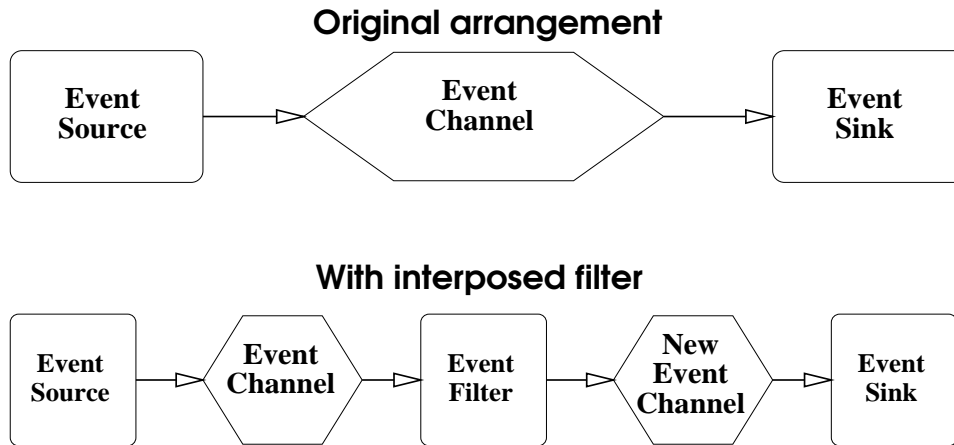


Figure 11: Source and sink with interposed event filter.

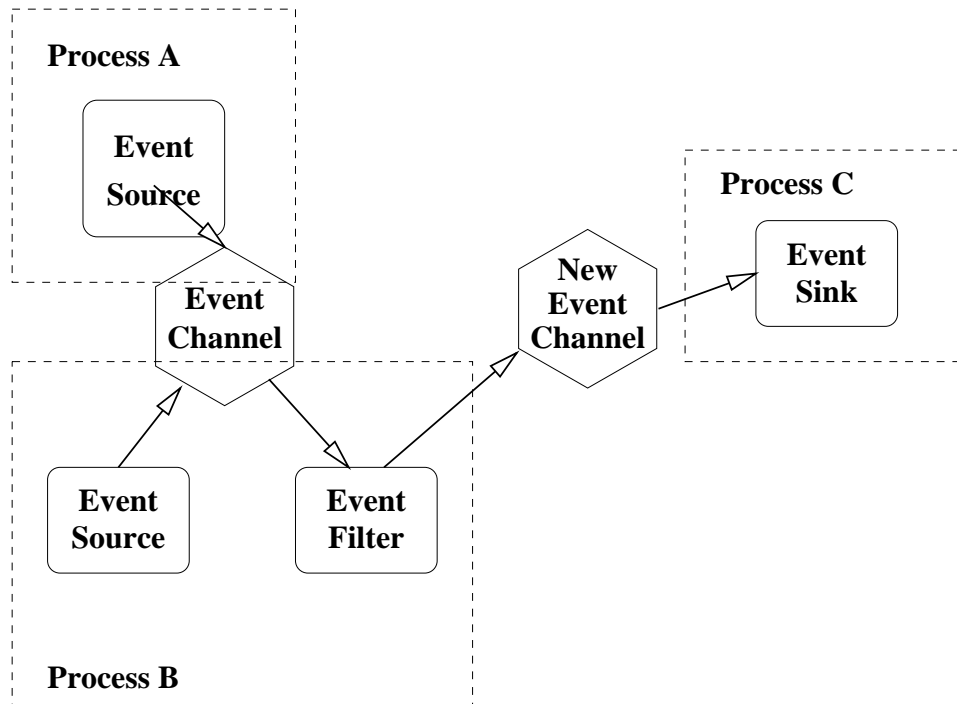


Figure 12: Filter with more than one source.

are submitted and transmit any event that results in the derived channel. This approach has the advantage that we limit unwanted event traffic (and the associated waste of compute and network resources) as much as possible. any of the sources in the original traffic, network traffic between those elements is avoided entirely. Figure 13 shows the logical arrangement of a derived event channel.

## 5.2 Mobile Functions and the E-code Language

A critical issue in the implementation of derived event channels is the nature of the function  $F$  and its specification. Since  $F$  is specified by the sink but must be evaluated at the (possibly remote) source, a simple function pointer is obviously insufficient. There are several possible approaches to this problem, including:

- severely restricting  $F$ , such as to preselected values or to boolean operators,
- relying on pre-generated shared object files, or
- using interpreted code.

Having a relatively restricted filter language, such as one limited to combinations of boolean operators, is the approach chosen in the CORBA Notification Services[7] and in Siena[1]. This approach facilitates efficient interpretation, but the restricted language may not be able to express the full range of conditions that may be useful to an application, thus limiting its applicability. To avoid this limitation it is desirable to express  $F$  in the form of a more general programming language. One might consider supplying  $F$  in the form of a shared object file that could be dynamically linked into the process of the event source. Using shared objects allows  $F$  to be a general function, but requires the sink to supply  $F$  as a native object file for each source. This is relatively easy in a

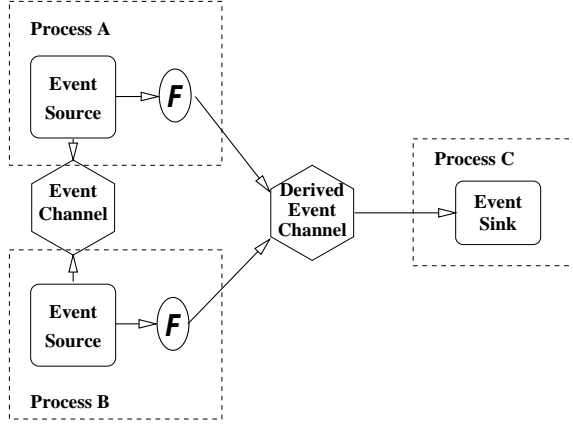


Figure 13: A derived event channel and function  $F$  moved to event sources.

homogeneous system, but becomes increasingly difficult as heterogeneity is introduced.

In order to avoid problems with heterogeneity one might supply  $F$  in an interpreted language, such as a TCL function or Java code. This would allow general functions and alleviate the difficulties with heterogeneity, but it impacts efficiency and requires a potentially large interpreter environment everywhere event channels are used. Given that many useful filter functions are quite simple and given our intended application in the area of high performance computing we rejected these approaches as unsuitable. Instead, we consider these and other approaches as a complement to the methods described next.

The approach taken in ECho preserves the expressiveness of a general programming language and the efficiency of shared objects while retaining the generality of interpreted languages. The function  $F$  is expressed in E-Code, a subset of a general language, and dynamic code generation is used to create a native version of  $F$  on the source host. E-Code may be extended as future needs warrant, but currently it is a subset of C. Currently it supports the C operators, **for** loops, **if** statements and **return** statements. Extensions to other language features are straightforward and several are under consideration as described in Section 6.

E-Code's dynamic code generation capabilities are based on Icode, an internal interface developed at MIT as part of the 'C project[10]. Icode is itself based on Vcode[5] also developed at MIT by Dawson Engler. Vcode supports dynamic code generation for MIPS, Alpha and Sparc processors. We have extended it to support MIPS n32 and 64-bit ABIs and x86 processors<sup>6</sup>. Vcode offers a virtual RISC instruction set for dynamic code generation. The Icode layer adds register allocation and assignment. E-Code consists primarily of a lexer, parser, semanticizer and code generator.

ECho currently supports derived event channels that use E-Code in two ways. In the first, the event type in the derived channel is the same as that of the channel from which it is derived (the parent channel). In this case, the E-Code required is a boolean filter function accepting a single parameter, the input event. If the function returns non-zero it is submitted to the derived event channel, otherwise it is filtered out. Event filters may be quite simple, such as the example below:

<sup>6</sup>Integer x86 support was developed at MIT. We extended Vcode to support the x86 floating point instruction set (only when used with Icode).

```
extern EChannel EChannel_derive(DExchange de, char *chan_id, char *filter_function);
extern EChannel EChannel_typed_derive(DExchange de, char *chan_id, char *filter_function,
                                      IOFieldList field_list, DEFormatList format_list);
```

Figure 14: Derived event channel API.

```
{
    if (input.level > 0.5) {
        return 1; /* submit event into derived channel */
    }
}
```

When used to derive a channel, this code is transported in string form to the event sources associated with the parent channel, is parsed and native code is generated at those points. The implicit context in which this code evaluated is a function declaration of the form:

```
int f(<input event type> input)
```

where *<input event type>* is the type associated with the parent channel.<sup>7</sup> The API for channel derivation is shown in Figure 14. Once derived, the created channel behaves as a normal channel with respect to sinks. It has all of the sources of the parent channel as implicit sources, but new sources providing unfiltered events can also be associated with it.

While this basic support for event filtering is a very powerful mechanism for suppressing unnecessary events in a distributed environment, ECho also supports derived event channels where the event types associated with the derived channel is not the same as that of the parent channel. In this case, E-Code is evaluated in the context of a function declaration of the form:

```
int f(<input event type> input, <output event type> output)
```

The return value continues to specify whether or not the event is to be submitted into the derived channel, but the differentiation between input and output events allows a new range of processing to be migrated to event sources.

One use for this capability is remote data reduction. For example, consider event channels used for monitoring of scientific calculations, such as the global climate model described in [8]. Further, consider a sink that may be interested in some property of the monitored data, such as an average value over the range of data. Instead of requiring the sink to receive the entire event and do its own data reduction we could save considerable network resources by just sending the average instead of the entire event data. This can be accomplished by deriving a channel using a function which performs the appropriate data reduction. For example, the following E-Code function:

```
{
    int i;
    int j;
    double sum = 0.0;
    for(i = 0; i<37; i= i+1) {
        for(j = 0; j<253; j=j+1) {
            sum = sum + input.wind_velocity[j][i];
        }
    }
    output.average_velocity = sum / (37 * 253);
    return 1;
}
```

---

<sup>7</sup>Since event types are required, new channels can only be derived from typed channels.

performs such an average over atmospheric data generated by the atmospheric simulation described in [8], reducing the amount of data to be transmitted by nearly four orders of magnitude.

## 6 Ongoing Work

Derived event channels are a recent result and are subject to continuing improvement as we gain experience with their use in real applications. However, some areas for future work are already apparent:

**Parameterized Filters** Currently derivation functions are simple functions of their input events. However, there are some obvious ways where more powerful functions could be valuable. Consider the situation where a sink wants a filter function based on values which change (hopefully more slowly than the event stream they are filtering). A simple example might occur in a distributed virtual reality application using event channels to share participant location information. Rather than sending location information as fast as the network allows, a more intelligent system might use derived event channels to turn down the update rate when the participants are not in sight of each other or merely distant. However, these conditions obviously change over time. One could periodically destroy and re-derive channels with updated filter functions, but a more straightforward approach would be to associate some state with a derivation function and allow it to be updated by the originator.

**Visibility into Source Process** [4] anticipates using the functionality offered by derived event channels for program steering. Currently external program steering typically requires monitoring events to be delivered to an external process. There they are interpreted and actions can be taken to affect the state of the application being steered. Like filter functions, steering decision functions may be as simple as comparing an incoming value with a threshold and steering the program if the threshold is exceeded. As a result, program steering requires a minimum of a network round-trip, and it is an asynchronous process. If the steering functions could be migrated into the steered application, much like we are moving filter functions to event sources, steering latencies could be reduced by several orders of magnitude and undertaken synchronously. When program steering is performed by a human-in-the-loop system, latency and synchronicity are of little importance. But steering can also be performed algorithmically and holds significant potential in terms of allowing programs, system abstractions and even operating systems to adapt to changing usage and resource situations[11, 9]. Such adaptations must be rapid and often must be synchronous in order to be valid.

However, doing program steering in a derivation function requires that the function have visibility into application program to call functions or affect values there. How this can be cleanly accomplished is an open question, but we anticipate mechanisms for associating a “context” with an event channel. The context would specify program elements which might be accessible to derivation functions. In this way, visibility into the sending application would be controlled by the sender but the use and updating of visible elements would be specified by the receiver through derivation function.

**Function Analysis** There is also potential benefit in specializing and analyzing the derivation functions. In particular, an event source clearly has an interest in understanding the nature of

the code it agrees to execute for a derivation function. While the simplest functions may be safe in the sense that they have no visibility into source's environment, they may walk off the end of event arrays or contain infinite loops. However, in this environment we know the size and extent of all data types. We can easily generate array bounds checks and can consider analyzing the specified functions for computational complexity and termination. An event source on a highly loaded machine might be given the option of rejecting a derivation that would increase its workload. More generally, since derived event channels always represent some tradeoff between compute and network resources, we can consider creating a more generalized mechanism through which tradeoff decisions can be made in the context of dynamic feedback from resource monitoring. Such a system could greatly ease the process of building applications which perform robustly in today's dynamic network environments.

**Optimization** The dynamic code generation tools used in Derived Event Channels generate code quickly, but the generated code is not well optimized. For example, the filter whose dynamic code generation was described above to have resulted in 21 Sparc instructions can be correctly performed in 10 instructions. Vcode includes a peephole optimizer for the Sparc architecture, but the generated code could clearly benefit from increased optimization. However, optimization is a significant source of complexity and compile-time cost in modern compilers. Whether the efficiency gains of additional optimization are significant enough to outweigh their increased costs is an avenue for future research.

## References

- [1] Antonio Carzaniga, David S. Rosenblum, and Alexander L. Wolf. Design of a scalable event notification service: Interface and architecture. Technical report, Department of Computer Science, University of Colorado at Boulder, August 1998. Technical Report CU-CS-863-98.
- [2] Greg Eisenhauer. Portable self-describing binary data streams. Technical Report GIT-CC-94-45, College of Computing, Georgia Institute of Technology, 1994. (*anon. ftp from ftp.cc.gatech.edu*).
- [3] Greg Eisenhauer, Beth Schroeder, and Karsten Schwan. Dataexchange: High performance communication in distributed laboratories. *Journal of Parallel Computing*, (24), 1998.
- [4] Greg Eisenhauer and Karsten Schwan. An object-based infrastructure for program monitoring and steering. In *Proceedings of the 2nd SIGMETRICS Symposium on Parallel and Distributed Tools (SPDT'98)*, pages 10–20, August 1998.
- [5] Dawson R. Engler. Vcode: a retargetable, extensible, very fast dynamic code generation system. In *Proceedings of ACM SIGPLAN'96 Conference on Programming Language Design and Implementation*, 1996.
- [6] Object Management Group. *CORBAservices: Common Object Services Specification*, chapter 4. OMG, 1997. <http://www.omg.org>.
- [7] Object Management Group. Notification service. <http://www.omg.org>, Document telecom/98-01-01, 1998.

- [8] Thomas Kindler, Karsten Schwan, Dilma Silva, Mary Trauner, and Fred Alyea. A parallel spectral model for atmospheric transport processes. *Concurrency: Practice and Experience*, 8(9):639–666, November 1996.
- [9] B. Mukherjee and K. Schwan. Implementation of scalable blocking locks using an adaptive threads scheduler. In *International Parallel Processing Symposium (IPPS)*. IEEE, April 1996.
- [10] Massimiliano Poletto, Dawson Engler, and M. Frans Kaashoek. tcc: A template-based compiler for ‘c. In *Proceedings of the First Workshop on Compiler Support for Systems Software (WCSS)*, February 1996.
- [11] Daniela Ivan Rosu, Karsten Schwan, Sudhakar Yalamanchili, and Rakesh Jha. On adaptive resource allocation for complex real-time applications. In *18th IEEE Real-Time Systems Symposium, San Francisco, CA*, pages 320–329. IEEE, Dec. 1997.